Artificial Intelligence 48 (1991) 57–86
Elsevier

Animate vision*

Dana H. Ballard

Computer Science Department, University of Rochester, Rochester, NY 14627, USA

Received March 1990
Revised June 1990

Abstract

Ballard, D.H., Animate vision, Artificial Intelligence 48 (1991) 57–86.

Animate vision systems have gaze control mechanisms that can actively position the camera coordinate system in response to physical stimuli. Compared to passive systems, animate systems show that visual computation can be vastly less expensive when considered in the larger context of behavior. The most important visual behavior is the ability to control the direction of gaze. This allows the use of very low resolution imaging that has a high virtual resolution. Using such a system in a controlled way provides additional constraints that dramatically simplify the computations of early vision. Another important behavior is the way the environment "behaves". Animate systems under real-time constraints can further reduce their computational burden by using environmental cues that are perspicuous in the local context. A third source of economy is introduced when behaviors are learned. Because errors are rarely fatal, systems using learning algorithms can amortize computational cost over extended periods. Further economies can be achieved when the learning system uses indexical reference, which is a form of dynamic variable binding. Animate vision is a natural way of implementing this dynamic binding.

1. What is vision for?

We are accustomed to thinking of the task of vision as being the construction of a detailed representation of the physical world. Furthermore, this constructive process is regarded as being independent of larger tasks. From the *Encyclopedia of Artificial Intelligence*: "the goal of an image understanding system is to transform two dimensional data into a description of the three dimensional spatiotemporal world" and such a system "must infer 3-D surfaces, volumes, boundaries, shadows, occlusion, depth, color, motion" [58, p. 389]. However, a paradigm that we term *animate vision*[1] argues that vision is

* Revised version of the paper that won the Artificial Intelligence Journal Best Paper Award at IJCAI-89, Detroit, MI.

This research was supported by NSF Grant No. DCR-8602958 and NIH Grant No. R01 NS22407-01.

[1] Why pick the term *animate vision* when there already is the notion of *active vision*? One problem with *active vision* is that it is readily confused with *active sensing*, which has been used for laser rangefinders, etc. Also it has been associated with multi-modal fusion [2] regardless of goals.

more readily understood in the context of the visual behaviors that the system is engaged in, and that these behaviors may not require elaborate categorical representations of the 3-D world. Animate visual systems have anthropomorphic features such as binocularity, foveas, and most importantly high speed gaze control. While it is possible to build many different kinds of visual systems, such as those that have more than two cameras or use active sensing, what we are calling animate vision is directed towards specific computational advantages of having anthropomorphic features. The main purpose of this paper is to summarize these computational advantages.

Throughout the paper we stress that whatever models are produced must function in real time. As a research stratagem, we shun general-purpose algorithms if they must appeal to vast increases in computing power in order to be practical. Instead our method is to look at ways to increase computational speed that exploit additional constraints introduced when the animate system is allowed to interact with its environment.

The goal of animate vision is the use of vision in behaviors associated with intelligence, and as such it has its roots in theories of robot behaviors. Brooks has argued for behaviors that do not require internal representations in a larger context [12, 13], and others have demonstrated the importance of active vision systems that integrate vision with behavior (Moravec [38], Bajcsy and Allen [4], Chen and Kak [17]) as well as demonstrating the advantages of knowing camera motions (Aloimonos et al. [2]). Ullman has emphasized the use of task-directed programs that operate on the optic array [62]. Animate vision also has its roots in the study of vision of the lower animals. From studies of the frog, Arbib [3] has long been stressing the integral role of vision in behavior as a *perception-action cycle*. Many of the technical features of insect vision can be used by animate vision systems and some of these have recently been realized by Nelson [40]. However, our primary purpose is to develop the advantages of animate vision that are geared towards hand-eye coordination behaviors. (Although this paper heavily emphasizes the role of the visual system and treats the hand only to the extent needed to explore some interactions.)

To start to see how animate vision might be qualitatively different from passive vision, let us examine the structure and function of eye movements in the human visual system. The human eye is distinguished from current electronic cameras by virtue of having much better resolution in a small region near the optical axis. This region is termed the fovea, and has a diameter approximately one to two degrees of visual angle. Over this region the resolution is better by an order of magnitude than that in the periphery. One feature of this design is the simultaneous representation of a large field of view and local high acuity. Figure 1, from a study by Sandini and Tagliasco [57], shows graphically the kind of gains that can be achieved.

Figure 1 visually understates the situation for the human system, where the

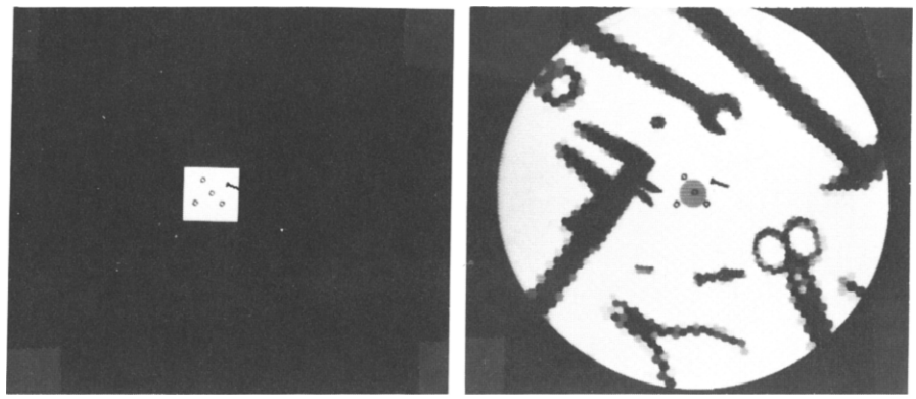

Fig. 1. (a) 700×700 image taken of a backlighted scene of industrial parts. At the given resolution, the field of view is very small. (b) The same number of samples using a logarithmic decrease in resolution from the optical axis. A dramatic increase in field of view is achieved at the price of peripheral resolution (Sandini and Tagliasco [57]).

fovea is less than 0.01% of the visual field area! With the small fovea at a premium in a large visual field, it is not surprising that the human visual system has special behaviors (saccades) for quickly moving the fovea to different spatial targets [42]. The first systematic study of saccadic eye movements in the context of behavior was done by Yarbus [68]. A selection of his data are shown in Fig. 2. Subjects were given specific tasks pertaining to a familiar picture. The figure shows the traces for three minutes of viewing as a subject attempts to solve different tasks: (a) give the ages of the people; (b) surmise what the family had been doing before the arrival of the "unexpected visitor"; and (c) remember the position of the people and the objects in the room. This data shows what has been confirmed by several other studies: Subjects use scanning patterns that are highly sensitive to the particular task at hand [43–45]. Of the traces in Fig. 2, the last is most remarkable, since it is so similar to the task of so many computer vision programs: we conjecture that since the eye movement traces show a specialized signature for this task as well, it is *not* done routinely. Instead, the overall impression of these traces is that the visual system is used to subserve problem-solving behaviors and such behaviors often do *not* require an accurate model of the world in the traditional sense of remembering positions of people and objects in a room.

The above data on the fovea and saccades hint also at how dynamic a process visual behavior must be. Saccades at the rate of three per second are routine in visual problem solving. Furthermore most of the brain structures that represent visual information are retinally indexed. This means that their state is changed with each eye movement. This raises a technical puzzle for human visual perception: How can the world appear to be stable when the data collecting process is so dynamic? We believe that this is a profound question with a surprising answer: The visual system provides the illusion of three-

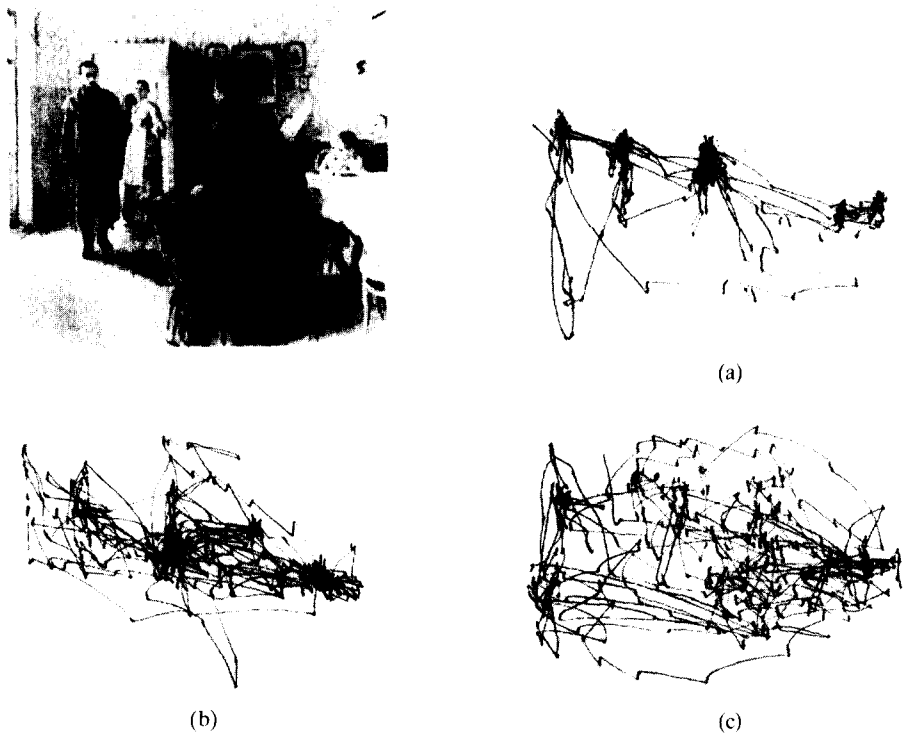

Fig. 2. (after [68]). Reproduction from I.E. Repin's picture "An Unexpected Visitor" and three records of eye movements. The subject examined the reproduction with both eyes for three minutes each time. Before the recording sessions, the subject was asked to: (a) give the ages of the people; (b) surmise what the family had been doing before the arrival of the unexpected visitor; and (c) remember the position of the people and objects in the room.

dimensional stability by virtue of being able to execute fast behaviors. This point may be very difficult as it is so counter-intuitive, but it has been arrived at in different forms by many different researchers. For example Rosenschein has stressed the importance of implicit knowledge representation by a behaving "situated automaton" [54, 55]. This may have been the point of Gibson's "affordances" [24]. O'Regan and Lévy-Schoen emphasize the use of the world as a "memory buffer" that can be accessed by visual behaviors [48]. Dickmanns's self-driven car makes extensive use of a dynamic model of the roadway [20]. At any rate, having a particular embodiment forces one to deal with performance issues: One has to act in a timely manner under resource constraints. One way to do this would be to have an elaborate internal representation as a form of "table look-up." But in a dynamic world, the cost of maintaining the correspondence between the representation and the world becomes prohibitive. For this reason animate vision systems may have to travel light and depend on highly adaptive behaviors that can quickly discover how to use current context.

We develop these ideas at three different levels of abstraction. Section 3 summarizes the computational advantages of an anthropomorphic gaze control, with particular emphasis on "early vision." In particular, we show how a particular anthropomorphic feature, the fixation frame, vastly simplifies the computation of physical invariances from photometric data. Section 4 shows how such computations can be integrated into complete behaviors such as searching for an object and recognizing an object. Section 4 also proposes a model of local spatial memory by showing how an animate agent can use knowledge of its recent history and extra-visual sensors to define geometric aspects of its environment. Finally, Section 5 shows how abstractions of these kinds of behaviors can be used in learning algorithms. The dynamic nature of the learning algorithms can further reduce the need for elaborate internal representations.

2. The animate vision paradigm

The central asset of animate vision is gaze control. Gaze control is the collection of different mechanisms for keeping the fovea over a given spatial target. The single most distinguishing feature of the human visual system is its high-speed gaze control mechanisms. As animals, we move in relatively fixed environments, but we also have to deal with other moving objects, animate and inanimate. Although we must function in the presence of different kinds of motion, our visual system works best when the imaged part of the world does not move. However, for a variety of behaviors, such as running after moving objects and hand–eye coordination, the complete visual field cannot be stabilized. Instead, stabilization can be achieved for a region near a point in the world near the optical axes that commands the viewer's gaze.[2] That point is termed the point of fixation and is defined by the intersection of the two optical axes.

Gaze control mechanisms fundamentally change computational models of vision. Without them the visual system must work in isolation, with the burden of solving difficult problems with many degrees of freedom. With them a new paradigm emerges in which the visual calculations are embedded in a sensory-motor behavioral repertoire. Rather than thinking of visual processing as separate from cognitive or motor processing, they are interlinked in terms of integral behaviors. These behaviors need not always be successful but they must be timely: Some competence may be sacrificed for timely performance. This viewpoint has many different kinds of advantages.

[2] Here we are neglecting the very small motions of the eye [42, p. 95] as unimportant in a behavioral context.

(1) *Animate vision systems can use physical search.* The system can move the cameras in order to get closer to objects, change focus, or change the point of view [29, 49, 65]. Often this visual search is more effective and less costly than algorithmic search on a single image, which may not even have the desired object in its field of view [41].

(2) *Animate vision can make (approximately) known camera movements.* Since these movements are self-generated, they provide additional constraints on the imaging process [2]. This facilitates the computational process dramatically: properties that are difficult to compute with a fixed camera system are much more easily computed with a moving camera system. One of the first demonstrations of this advantage was Bandopadhay's computation of rigid body motion parameters [8].

(3) *Animate vision can use exocentric coordinate frames.* The ability to control the camera's gaze, particularly the ability to fixate targets in the world while in motion, allows a robot to choose external coordinate frames that are attached to points in the world (see Fig. 3). Behaviors based on fixation point relative coordinates allow visual computations to be done with less precision.

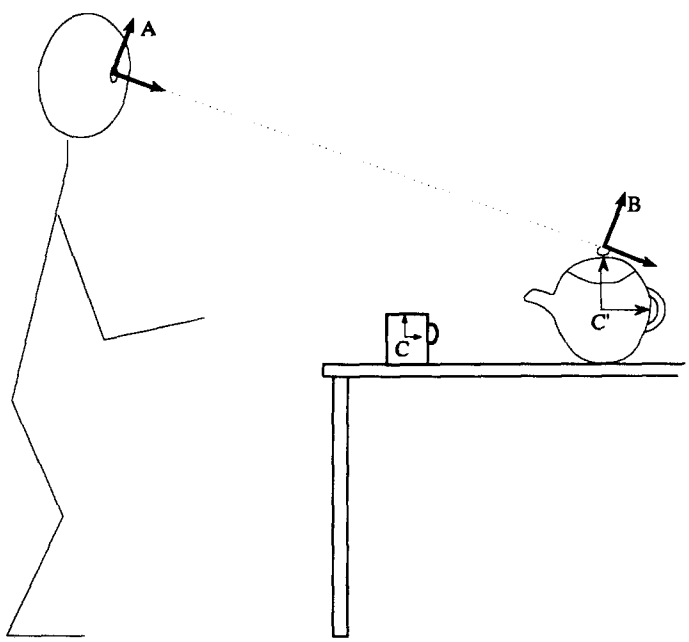

Fig. 3. Much previous work in computational vision has assumed that the vision system is passive and computations are performed in a viewer-centered frame (A). Instead, biological and psychophysical data argue for a world-centered frame (B). This frame is selected by the observer to suit information-gathering goals and is centered at the fixation point. The task of the observer is to relate information in the fixation point frame to object-centered frames (C).

(4) *Animate vision can use relative (or qualitative) algorithms.* The fixation point reference frame allows visuo-motor control strategies that servo relative to that frame. These are much simpler than strategies that use egocentric coordinates.

(5) *Gaze control can segment areas of interest in the image precategorically.* That is, one can isolate candidate visual features without first associating them with models using the degrees of freedom of the gaze control mechanisms. For example, one can use the blurring introduced by self-motion while fixating to isolate the region around the point of fixation [16]. Similarly, one can use regions of near zero disparity produced by a binocular vergence system.

(6) *Animate systems can exploit environmental context.* Gaze control leads naturally to the use of object-centered coordinate systems as the basis for spatial memory. Object-centered coordinates have a great advantage over egocentric coordinates in that they are invariant with respect to observer motion. Keeping track of relations between object-centered frames allows for simplified object location strategies.

(7) *Animate vision is tailor-made for learning algorithms that use indexical reference.* Gaze control with a high resolution fovea to isolate visual features is tailor-made for systems that use indexical reference [1, 64]. Such systems provide a controlled access to the environment, making it much easier to access stored plans. Furthermore, such systems are tailor-made for reinforcement learning algorithms. The vast reduction in the state space provided by indexical reference makes the use of such brute force learning algorithms possible. In turn, learning algorithms allow visual behaviors to learn just those features that are useful for solving the problem in very specific contexts. This leads to further computational economies.

3. The fixation frame

One of the most central aspects of animate vision is the use of an exocentric coordinate frame termed the frame of fixation. This frame provides direct access to information from a small region near the fixated point. Of particular importance is the information associated with *early vision* [33]. Early vision builds retinotopically indexed maps of important environmental features such as depth, color, and velocity. Despite extensive work in this area over the past decade, the construction of such maps with computational models has proven to be very difficult. A primary reason for this may have been the assumption of a passive vision system. In an animate vision system, the degrees of freedom of the cameras are under the control of the animal. Aloimonos et al. [2] show in a general way how such assumptions can stabilize the computation of those

features but their analysis misses the following vital point. A passive vision system is more or less constrained to use the coordinate system dictated by the camera optics. In contrast, an active system that can fixate an environmental point can use an object-centered frame of reference centered at that point. The calculations of early vision are greatly simplified given this ability. Note that this is a very different assertion than that of Marr [33], who emphasized that the calculations were in viewer-centered coordinates. We assert that the calculations are more correctly represented as being in world-centered coordinates. As shown in Fig. 3, the world-centered frame is viewer-oriented, but not viewer-centered. The origin of this frame is at the point of intersection of the two optical axes. To orient this frame one axis can be parallel to the line joining the two camera centers; the other can be chosen as the optical axis of the dominant eye.[3]

3.1. *Using the fixation frame*

To illustrate the advantages of using the fixation frame, we developed a computational model of motion parallax. Motion parallax, or kinetic depth, is the sensation of depth obtained by moving the head while fixating an environmental point in a static scene. If the observer has little forward motion, objects in front of the fixation point appear to move in the opposite direction to the motion while objects behind the fixation point move in the same direction. (For a more general analysis that includes forward motion, see [52].) The apparent velocity is proportional to the distance from the fixation point [19]. Under these conditions it is easy to compute scaled depth (depth/fixation depth), which is a monotonic function of spatial and temporal derivatives of the image intensity function and has a zero value at the fixation point. By implementing this strategy on our robot we verified that a depth estimate can be obtained in real time over a 400×400 pixel image without iteration [7].[4] This result shows that the early vision computations of animate vision, at least in the case of kinetic depth, are decidedly simpler than fixed camera vision, as first noted by Aloimonos et al. [2]. Table 1 compares the two paradigms.

3.2. *Gaze control*

The small size of the fovea, together with the rapid movements humans can make, places a premium on gaze stabilization mechanisms. Perhaps for this

[3] This can be a subtle distinction, especially since animal data show that visual information in the cortex is retinotopically indexed. However, the distinguishing feature is the logical zero of the coordinate system: For a system with gaze control, zero velocity and zero disparity are located at the fixation point.

[4] The local nature of the computations make them ideal for implementation by pipeline computer architectures. Such architectures pass the digitized signals at video frame rates through a succession of special function processors. The modularity of these architectures, together with their video frame rate speed, is revolutionizing real-time image processing.

Table 1
A comparison of the computational features of fixed camera vision and animate vision.

Fixed camera vision	Animate vision
Local constraints that relate physical parameters to photometric parameters are underdetermined.	Local constraints are sufficient.
Minimalist constraints such as smoothness used to regularize the solution.	Maximalist constraints such as specific behavioral assumptions used to obtain solution
Algorithm requires parallel iterations over the retinally indexed array.	Algorithm is local and has a constant time solution.
Frame of reference is camera-centered (egocentric).	Frame of reference is fixation point centered (exocentric).

reason a number of separate mechanisms for human gaze control have evolved. As Table 2 shows, the eye movement system has a number of different systems that function to control gaze under different circumstances. In addition there is the accommodation system that acts to focus the lens.

We argue that the ability to control gaze can greatly simplify the computations of early vision, but what of the complexity of gaze control itself? If that should turn out to be prohibitively difficult it would negate the value of this paradigm. Fortunately, all our experimental work to date argues that this will not be the case [6], as does work by Clark and Ferrier [18]. Figure 4 shows our animate vision system. Currently we use a "dominant eye" control protocol whereby the dominant camera controls the system pitch and its own yaw coordinate using a simple correlation tracking scheme [16]. The non-dominant camera uses a novel vergence correction algorithm [47] based on the cepstral filter [69] to correct its own yaw error. Brown [14, 15, 53] has recently shown how these and other components can work together synergistically. These components run in real time. At the moment there are many differences with a reasonable human model, but the performance is sufficiently good to allow us to explore vision while fixating in real time. Details may be found in [16].

Table 2
Summary of primate gaze control systems.

Hold gaze	Fixed target	Vestibular-ocular reflex (VOR). A system that uses knowledge of accommodation and vergence state together with head accelerations to stabilize the gaze vector.
	Moving target	Vergence. A binocular system for locking both foveas over the same three-dimensional target.
		Pursuit. A system for tracking moving objects by generating smooth velocity control signals.
Change gaze	Saccades. High speed precomputed movements that rapidly change gaze over small to very large visual angles.	

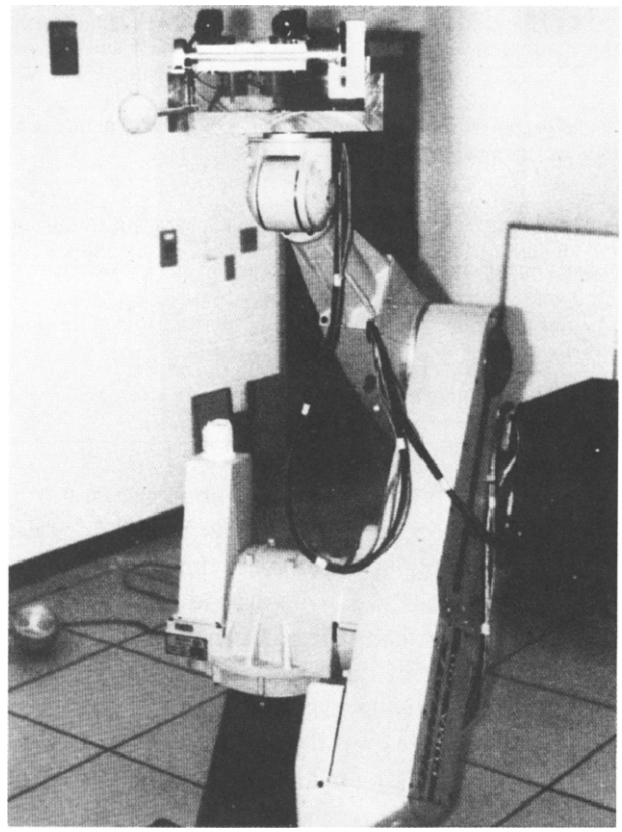

Fig. 4. The University of Rochester's animate vision system. The "robot head" has three motors and two CCD high-resolution television cameras providing input to a DataCube MaxVideo® image-processing system. One motor controls pitch of the two-eye platform, and separate motors control each camera's yaw. The motors have a resolution of 2,500 positions per revolution and a maximum speed of 400°/second. The robot arm, a Unimation 762, has a workspace consisting of most of the volume of a sphere with a two-meter radius, and a top speed of about one meter per second. The first such system, built at the University of Pennsylvania by Bajcsy [4], demonstrated the potential for vision with controlled cameras. It had vergence and accommodation and zoom control. The main drawbacks were its slow speed and limited workspace.

The importance of vergence in gaze control is dramatically demonstrated by Olson and Potter [47]. Without vergence, very large disparities on the order of half the image dimension can be obtained. These pose difficulties for algorithms that use stereo to build depth maps. With vergence, the disparities for the objects of interest can be kept small. In fact, most models of human stereopsis posit or require a fusional system that brings the disparities within the range of a detailed correspondence process [21, 33, 69].

3.3. Relative vision

The kinetic depth computation naturally produces a relative result; for absolute depth the calculations must include direction of gaze. Since the

relative result is easier to obtain, it motivates the question as to whether other visual behaviors might in fact use relative vision. In fact many examples can be found that suggest that relative quantities are used and that their computation is simpler. For example, many psychophysical tasks suggest that the way the image is interpreted depends on occlusion cues such as shown in Fig. 5 [39]. It is not easy to make such judgements from an arbitrary viewing position, as would be required by a viewer-centered hypothesis. The kinetic depth result suggests that the notion of a fixation point may be implicit behind the analysis even though we might not be aware of it. Our perceptual system is structured to make accurate judgements relative to an object-centered frame at the fixation depth. Simplistically, imagine that one keeps two maps: one for structures that are judged to be in front of or at the fixation depth, and one for structures that are behind the fixation depth. The different interpolation rules can be fixed for each map. This structure is much simpler than that which would be required for viewer-centered maps. Such maps would have to be able to make corrections based on comparisons of all possible pairs of depth values.

The notion that the computational results of early vision are intrinsically relative can be challenged by obvious counter-examples. We can reach our

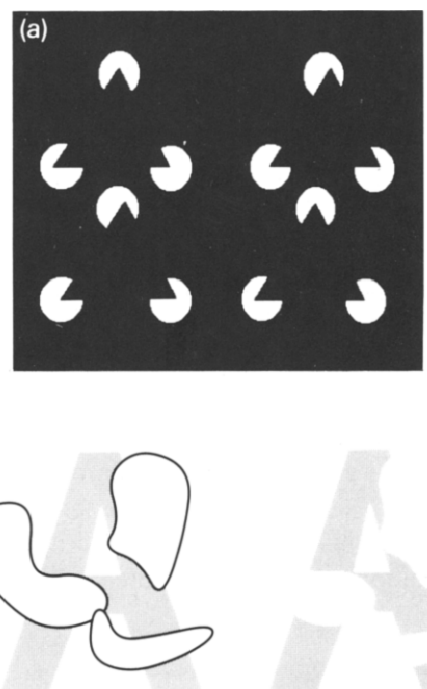

Fig. 5. (a) Ken Nakayama's [39] illusion of subjective contours using stereo (not to scale). When fused, if the relative disparities are such that the triangle is in front of the circles, subjective contours are seen; if behind, then they are not. (b) The letter "A" is easier to see if its components result from real occluding boundaries. This can be explained if the occluders can invoke a fixation depth that is in front of the plane of the A's components.

arms to places in space that we have recently seen but are not currently looking at. So at least in this example the information from the gaze system must be cast in three-dimensional coordinates. However, it may well be the case that the behavior that controls reaching under these circumstances is separate from behaviors that use relative visual data. One example of the latter comes from experiments done by Erkelens and Collewijn [21]. Subjects fixated a visual stimulus of random dot targets arranged in a fronto-parallel plane on a display. Changing the disparity of the dots created a situation similar to that which would have been produced by a real target moving back and forth in depth. In fact the subjects' vergence movements showed that they were tracking the simulated movement. In spite of this they reported no depth change in the perception of the plane. So even though the vergence state could in principle have been used to create a three-dimensional percept of a moving plane it was not done in this case. However, when another disparity was introduced into the display that remained fixed, subjects immediately saw the movement of the plane. The experiments of Olson and Potter [47] hint at why this might be the case. They used the central disparity target as a servo error signal to make the vergence system work. Owing to the small spatial extent of the binocular foveas, the vergence system must continually correct the gaze so that both foveas are verged on the same target. This means that the logical zero for this system is at the fixation point, in the same way that the logical zero for the kinetic depth system was at the fixation point.

The relative system has the virtue of requiring much less mathematical precision than the computations done in absolute coordinates. This is because the foveas provide the best precision only at the fixation point and an animate vision system can control the location of its fixation point. In contrast, to provide the same resolution everywhere, a fixed resolution system would have to be at least ten thousand times larger. This system would require even greater increases in computational costs, which scale by at least a low-order polynomial factor [61]. To see how the relative measurements could be used for three-dimensional positioning, consider visually guided reaching. An arm out of the plane of fixation can be guided in depth to a target at the fixation plane by using only relative disparities of the manipulator as seen by the visual system; the three-dimensional coordinates of the target are not required. This scheme also has the virtue of using the natural output of the stereo system which is in terms of fixation-relative coordinates.

4. Visual behaviors

A feature of the kinetic depth result is that it is an integral part of a visual behavior. When fixating a stationary point, the optical flow map can be interpreted as a depth map, but when pursuing a moving target, this interpreta-

tion is no longer valid. It could be the case that the kinetic depth result is an isolated case where behavior makes a large difference in the complexity of a problem in visual computation. However, a survey of the vision literature shows that there are many examples, including some very important recent cases, where the inclusion of behavior simplifies the computation. Behavior is used here in a very general sense to capture the self-motion of the animate system as well as the structure of the environment ("the behavior of the environment") in which the system operates. Table 3 summarizes some of these results.

If these special-purpose algorithms were the rule rather than the exception, then it may be that the visuo-motor system is best thought of as a very large amount of distinct special-purpose algorithms where the results of a computation can only be interpreted if the behavioral state is known. Ramachandran [51] has raised a similar point, arguing from psychophysical grounds that the visual system may best be thought of as many different algorithms that exploit different cues, but that do not always work and may not be simultaneously satisfiable. Brooks [12, 13] has also noted this point, using the term "sensor fission" to emphasize that different sensors may be used in different tasks. Recent work by Pentland on the shape from shading problem has also shown very simplified solutions for dominant special cases that depend on the behavioral milieu [50], and there have long been special case solutions to the motion problem that depend on behaviors. A compelling example of the central role of behavior in an animal system comes from Maunsell and Van Essen's work [35] on the macaque monkey. The macaque contains a very distinct retinotopic cortical map that is sensitive to motion. Regular electrode sampling across this map showed that the cortical visual area where the hands would be in hand–eye coordination, known as MT, was over-represented with

Table 3
Computations simplified by behavioral assumptions.

Agent's behavior	Behavioral assumption	References
Shape from shading	Light source not directly behind viewer.	Pentland [50]
Time to adjacency	Rectilinear motion; gaze in the direction of motion.	Lee and Lishman [70]
Kinetic depth	Lateral head motion while fixating a point in a stationary world.	Ballard and Ozcandarli [7]
Color homing	Target object is distinguished by its color spectrum.	Wixson and Ballard [65]
Edge homing	Target position can be described by approximate directions from texture in its surround.	Nelson and Aloimonos [41]

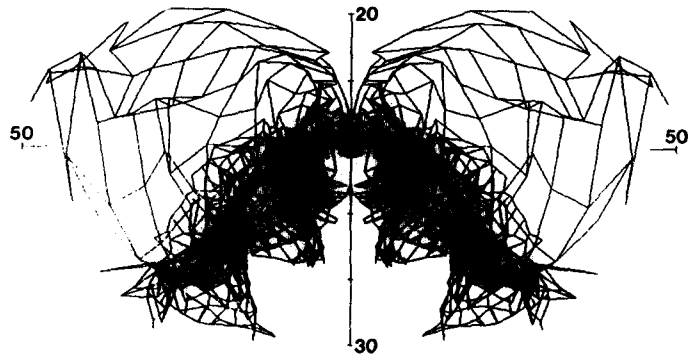

Fig. 6. Maunsell and Van Essen's [35] plot of the visual field in a macaque monkey's representation of optic flow shows that the portion of the visual field where the hands would be in a hand–eye coordination task is more densely represented than other areas. Numbers mark degrees. Data from one hemifield is reflected about the midline.

respect to other areas (Fig. 6). Experiments that record from a cortical area adjacent to area MT suggest specializations for behaviors that use foveal motion and behaviors that use peripheral motion [28].

In contrast to the notion of collections of behaviors, much of vision research has focused on a reductionistic approach whereby one tries to show how a particular quantity is computed independently of the behaviors that use it. Thus an opposing view of the results in Table 3 would be that they are too specialized and that general solutions should be sought. However, these general solutions have the price of an increased computational burden, and the demand for (timely) solutions in animate systems rules out all but relatively low complexity algorithms.[5]

One huge problem that remains is the real-time system problem of managing behaviors that compete for the imaging resources, and we have little to offer here. However, there are at least two situations where the competition is reduced: one where the processes have complementary activating conditions and another where they use different degrees of freedom of the motor resources. Visually-guided hand movements are sufficiently demanding behaviors that the advantages of a distributed approach to their control can be clearly demonstrated. In one of our laboratory experiments to test such a paradigm on a small scale, a robot system was developed to keep a balloon in the air by batting it with a paddle. The gaze control system was coupled to the motion of the robot via five completely independent visually-guided behaviors. Three of these controlled the position of the paddle (in height, width, and depth, respectively), one generated a batting movement, and another was responsible for re-acquiring the balloon visually when it escaped the field of view. There was no executive coordination of these systems at all, and, once

[5] Parallel architectures help but not with infeasible algorithms [61].

initialized, no communication between them except indirectly through their effects on the robot and the environment. Our preliminary experiments have revealed a remarkable level of coherent behavior using this strategy [67].

4.1. Quickly computable features

The real-time stress of animate vision requires that the kinds of visual cues used are easily computable. In a human system, the short fixation times are about 0.25 seconds and cortical neurons typically fire at rates of 10 spikes per second, leaving 2.5 spikes per fixation. In a sequential computer system using a 500×500 pixel image at video frame rates the demand is also great as the system must compute at roughly 10^7 pixels/second. Algorithms that require 10^3 instructions per pixel are common, leading to a demand of 10^{10} instructions/second.

One feature that is easily computed is color. Color has been neglected recently as a useful cue, although it has been used in earlier work (Feldman and Yakimovsky [22], Garvey [23], Beveridge et al. [9]). One reason for this neglect may have been the lack of good algorithms for color constancy. However, recently there has been great progress in correcting for both the chromaticity of the illuminant [32, 56] and for geometric effects such as specularity [27]. Another reason that color may not have been so successful is that it has been associated with a Mondrian-like view: one color per object. But many objects are multi-colored and this fact can prove very useful, as will be shown in the next section. A third reason for the neglect of color may be that it is not intrinsically related to the object's identity in the way that other cues, e.g., form, are. This view is well represented by Biederman [10]:

> Surface characteristics such as color and texture will typically have only secondary roles in primal access . . . we may know that a chair has a particular color and texture simultaneously with its volumetric description, but it is only the volumetric description that provides efficient access to the representation of CHAIR

but it is easily challenged. There are many examples from nature where color is used by animals and plants to send clear messages of enticement or warning. The manufacturing sector uses color extensively in packaging to market goods (e.g., Kodak). Animate vision systems can also use representations that are heavily *personalized* to achieve efficient behaviors, and color is an important feature for such representations. For example, it may not be helpful to model coffee cups as being red and white, but *mine* is, and that color combination is very useful in locating it. Another obvious example is commercial food packaging. We can readily describe the color of food packages for the kind of eggs and milk we buy even though these colors do not generalize: they will not work for another supermarket chain.

In summary, there have been various reasons for not using color, but most of these are now less compelling, particularly in the light of recent technical advances in color constancy and in reconsideration of the behavioral context in which color can be used. More importantly, color has two very important properties that make it a useful feature. Given that reasonable color constancy can be achieved, color has enormous value in vision as a cue because it is a punctate property of individual photoreceptors. This means that it is a very useful cue under conditions of low spatial resolution; precisely the conditions that exist in the periphery of the retina. The second useful property is view invariance. The colors of an object typically are invariant to wide ranges in field of view and to several different kinds of occlusion.

One way to take advantage of these properties uses the color histogram. Given a discrete color space, the color histogram is obtained by integrating over the image array:

$$h(\boldsymbol{c}) = \int f(\boldsymbol{c}, \boldsymbol{x})\, \mathrm{d}\boldsymbol{x} .$$

The color vector $\boldsymbol{c} = (r, g, b)$ obtained from the tri-chromatic receptor array can be sensitive to gross lighting changes such as the $1/r^2$ falloff from a point source. One way to compensate for this, observed in biological systems, is to use an opponent color space $\boldsymbol{c}' = (r - g,\ b - \frac{1}{2}(r + g))$. Figure 7 shows the

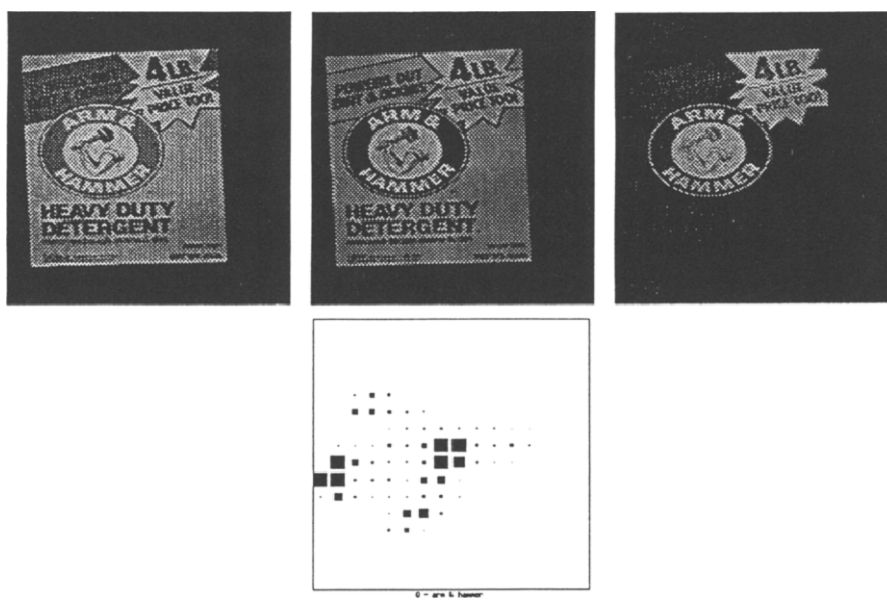

Fig. 7. Top: Red, green, and blue bands of "Arm & Hammer" image. The main body is yellow, the circle containing the hammer is red, the stripe at the top is green, and the lettering and hammer is blue and white. Bottom: Opponent color histogram of "Arm & Hammer" image, 16 buckets along each axis. Red–green axis runs vertically, green at the top, red at the bottom. Blue–yellow axis runs horizontally, yellow at the left, blue at the right. The yellow (far left) peak and black background (the center) peaks are the largest, and red and green peaks, as well as a small blue peak, are present. From Wixson and Ballard [65].

three chromatic channels from a color camera together with the opponent color histogram.

4.2. What/where behaviors

Returning to the challenge of Fig. 2(c), it seems that for human vision, locations and identities of objects are not routinely computed. Furthermore, the saccadic traces suggest that when this is done, the resultant computation requires many sequential eye movements. We further suggest that very different *algorithms* are used depending on the task of the moment. A gross distinction that can be made is between *identification algorithms* that analyze the foveated area during fixation and *location algorithms* that direct the eyes to new targets. Support for this WHAT/WHERE distinction, made by Mishkin [36, 37], comes from studies of human and primate brains. A major feature of the gross organization of the primate visual brain is the specialization of the temporal and parietal lobes of visual cortex [34, 36, 37]. The parietal cortex seems to be subserving the management of locations in space whereas the temporal cortex seems to be subserving the identification of objects in the case where location is not the issue. In a striking experiment by Mishkin [36], monkeys with parietal lesions fail at a task that requires using a relational cue but have no trouble performing a very similar task that requires using a pattern cue. The reverse is true for temporal lesions.

Why should the primate brain be specialized into two separate areas that are crucial for different functions? If we think generally about the problem of relating internal models to objects in the world, then one way to interpret this dichotomy is as a suggestion that the general problem of associating many models to many parts of the image simultaneously is too difficult. In order to make it computationally tractable within a single fixation, it has to be simplified, either into one of location (one internal model) or identification (one world object). Table 4 makes this suggestion more concrete.

Let us try to make the value of this dichotomy clearer through two specific examples, one involving a location behavior and one involving an identification

Table 4
The biological organization of cortex into WHAT/WHERE modules may have a basis in computational complexity. Trying to match a large number of image segments to a large number of models at once may be too difficult.

		Models	
		One	Many
Image parts	One	Manipulation. Trying to do some thing with an object whose identity and location are known.	Identification. Trying to identify an object whose location can be fixated.
	Many	Location. Trying to find a known object that may not be in view.	Too difficult?

behavior.[6] Both of these examples use the color histogram or spectogram as a central low-cost representation. This histogram can be used in two very different ways for different behaviors. If the location of a single known multicolored object is sought, the histogram of the current scene can be matched against that of the desired object. A robot can be trained to move toward the object by using this match function as a gradient. Let $M(h_m, h_i, \boldsymbol{x})$ be a function that scores the match between the object histogram and scene histogram at pose $\boldsymbol{x}$. For example, one possible match function is simply $\|h_m - h_i\|$. Now the robot can move in a direction of maximum $dM/d\boldsymbol{x}$. This works largely because different colors superpose in the color histogram, but if the *spectral* resolution is sufficient, they will not mix. The match function acts as a qualitative measure to direct the search. For an initial point one cannot depend on the object being within view, but if it is potentially viewable, an animate system can conduct a coarse saccadic scan of the view space and select good candidates by applying the match function to all of these discrete views. Figure 8 shows the results of doing this for two cases of looking for brightly colored objects.

Now let us turn to the complementary task: that of identifying an object whose location is known. The object can be isolated in various ways; one uses motion under fixation to blur nearby structure [46, 59]. If the image is assumed to be obtained from a single multicolored object, the histogram can be used as a multidimensional index into a database of multicolored objects. Given the notion of a match function $M(h_m, h_i)$, it is easy to find the model in terms of the best match, i.e.,

$$\left\{ m^* \,|\, M(h_{m^*}, h_i) = \max_m M(h_m, h_i) \right\}.$$

There are ways to perform this computation that are more efficient than a linear search through all the models, and details may be found in [60].[7] Figure 9 shows the results of matching nineteen objects one at a time into a database also of nineteen objects, but taken from different poses. Calculations by Swain show that the three-dimensional histogram has a very large capacity, given that the multicolored objects are distributed in color space.

These two examples of location and indexing are very simple when treated as separate behaviors, but would be difficult to combine into a single behavior or algorithm using many models and many image fragments (see the "too difficult" entry in Table 4). For example, trying to locate many objects simultaneously forces the different objects to compete for the peripheral

[6] We are not offering this as a proof that the brain uses color in this way. However, it is interesting that computational divisions suggested by brain architecture lead to vast simplifications in computation.

[7] These indexing experiments use the red–green–blue three-dimensional histogram instead of the opponent color histogram.

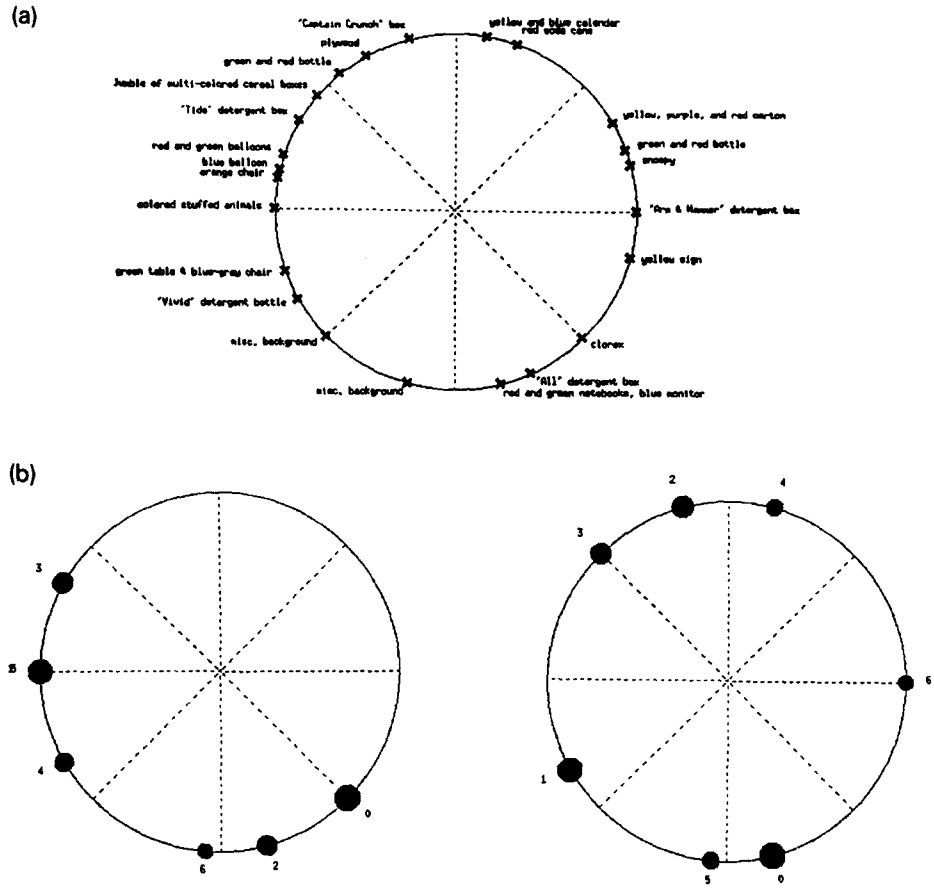

Fig. 8. Top view of the laboratory environment for a typical test run showing the direction (but not the distance) of each object with respect to the robot. The robot is in the center of the figure and common objects, denoted by filled circles, are located along the gaze directions shown. (b) Gaze directions produced by the object search mechanism for the "Clorox" and "All" detergent boxes. Area of circle is proportional to the confidence in that gaze. Numbers next to circles reflect the ordering of the confidences in decreasing order. From Wixson and Ballard [65].

resources of the animate system. Also, if many different models are placed into the model histogram h_m simultaneously, the effect of the cross-product of all the different colors is potentially devastating.

4.3. Spatial memory

The previous section explored one way of managing space, and that was homing. In the location task, a color signal was used to move the robot near a colored object. The homing behavior can be extended to a path using several landmarks, but each landmark must be in view at the appropriate time [41]. If the landmarks are not in view, the animate system has to resort to some kind of exhaustive search of visual space using its physical resources. Thus homing is

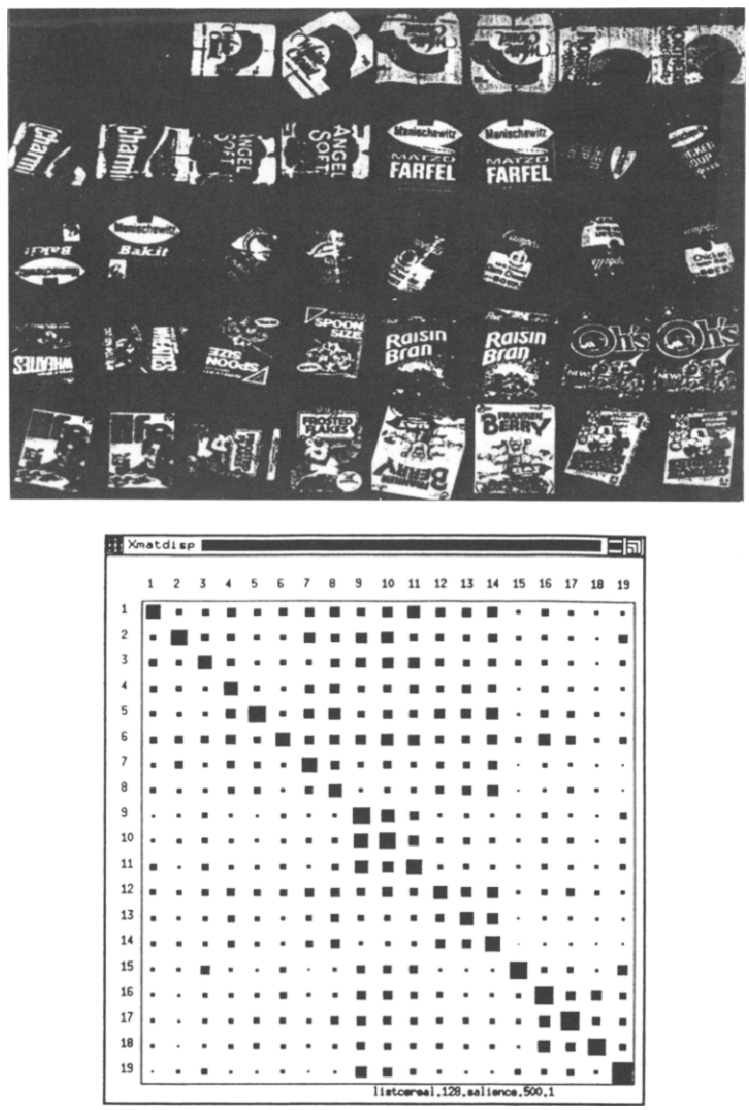

Fig. 9. Model indexing experiment based on color cues. Each of the nineteen models (upper image) is represented by its color histogram. Each of the unknown objects (to the right of each known object) is identified with the database color histogram that best matches its own color histogram. The results of matching all combinations of image and database histograms are displayed pictorially (lower image) where the sizes of the squares are proportional to match values. The dominance of the diagonal values shows that the correct match is always selected. From Swain and Ballard [60].

robust but can be expensive. To have more complicated behaviors than homing, some additional spatial memory structure is necessary. One extremist solution is to keep very high-resolution maps of the spatial environment and update these maps when something is changed. But for a variety of reasons, such a solution is not practical for animate systems. The foremost of these is

the errors in the measurement system itself, which are a function of the relative positions of the robot and target object. Another reason is that such maps are very expensive in terms of size, since only a small portion of the material is relevant to tasks that require it to be identified. A third reason is the expensive updating introduced by self-motion when the entire environment undergoes relative motion.

We have argued that animate vision allows the perception of properties of the world to be related to a coordinate frame that is attached to the world by using the abilities to fixate or pursue. However, this coordinate frame is only valid for the duration of the camera fixation; some additional structure is necessary for spatial memory. Thus for a variety of other reasons we need to introduce the notion of object-centered reference frames: (1) such frames allow the memory of objects' locations with respect to each other; (2) objects may be in motion; and (3) objects may not be in view. An elegant way of relating this coordinate frame to object-centered frames (OCFs) posits an explicit representation of transformations between OCFs and the current view. If one assumes that the model and view have primitive parts, for example, line segments, matches between these parts determine particular values of the transformation that relates the stored model to the current view [5, 25].

Figure 2 can be used to summarize the proposal for spatial memory. The current view represents similar features but with respect to a frame that is centered on the current fixation point (as opposed to the camera frame used by passive systems). For example, if the fixation point is the object-centered frame origin, the transformation will only differ by a rotation, having a translation value of zero. Spatial memory stores relationships between object-centered frames. In a computational theory of active vision, eye movements have an integral role in the storing and retrieval of spatial information in the following ways:

(1) The view transform T_{bc} contains the information necessary to foveate a visible object that has been recognized.
(2) Stored relationships between objects, $T_{cc'}$, can be used to transfer gaze from one object to another.

In contrast, egocentric or camera-centered systems attempt to maintain the transformations T_{ac} and $T_{ac'}$, which is more computationally intensive.

As noted in the introduction, the fovea is an elegant solution to the problem of simultaneously having high spatial resolution and a wide field of view given a fixed amount of imaging hardware. The price paid is that the target must be foveated. Thus small objects in a cluttered periphery can be effectively invisible. This means that directed visual search strategies must be employed to find objects. Think of car keys: to be useful, at any one time they must be kept in a familiar relationship with a large object. We think this difficulty can be minimized by having a stored model database whereby small objects are linked

to larger objects. To illustrate this proposal, we have built a two-dimensional eye movement simulator. Figure 10 shows the results from a test simulation. The problem is to locate a cup that is initially invisible in the periphery. Knowing that the cup is on the table, we first locate the table via a Hough transform technique [5] and then use the pose information to center the gaze. In this instance, once the gaze is centered on the table, the cup is within the high resolution fovea and can be found by using the same Hough transform technique, but now with the cup as the stored model. Here again, application of a system with a high precision fovea avoids the complexity of making fine-grained measurements over the full field of view.

Early work in vision attempted to use context in object recognition [23], but this work languished with the introduction of the Marr paradigm and its focus on low level vision. Since then, object recognition work has been very reluctant to use any kinds of context, with the result that object recognition is usually considered in a vacuum. The motivation for this is that general-purpose techniques that make few assumptions about the world would be more useful than special-purpose techniques. However, the disadvantage of this minimalist position is that methods with few assumptions typically fall back on search, which can lead to impractical computational demands. Instead of this minimalist notion of generality, animate vision advocates making maximal use of all the different kinds of constraints available. These are of two principal kinds.

(1) In a human or robot, one source of information is behavioral state. Humans have a vestibular system that measures linear and angular accelerations. This provides a short-term history of movements in the environment and also a measure of gravitational force. Another source is the human proprioceptive system, which provides the kinematic state as well as muscle torques.

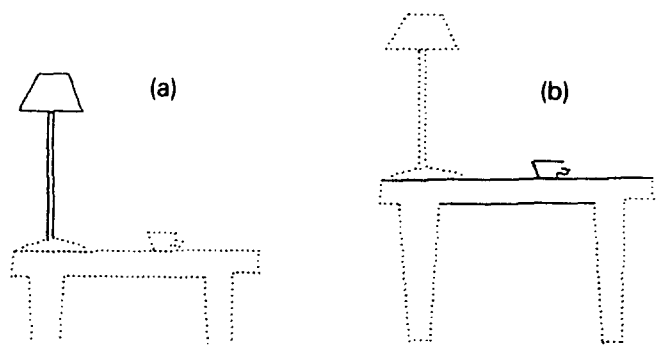

Fig. 10. A foveal vision system is an elegant solution to the problem of high spatial resolution and a wide field of view. The price paid is that small objects on the periphery are hard to see. However, known relationships with large objects can help. In (A), the cup cannot be easily seen, but in searching for the cup, one can first look for the table (B), which in this case brings the cup near the fovea, where it can be found.

(2) A second source of information is the local context in which objects appear. Objects are dependent on the surfaces of other objects for support. For example, chairs are supported by the ground plane, and pens are usually on tables. The design of objects in terms of support relationships constrains the way in which they interact with supporting surfaces. For example, chairs and cups usually have only three degrees of freedom, while in contact with their supporting surface: one rotation and two translation. If there were a way of exploiting these constraints it should make the recognition problem computationally simpler. The constraints supplied by behavioral state necessarily interact with those supplied by local context since, as animals, we have our own support needs. Thus we can use kinematics to directly measure the orientation of a ground plane or table surface with respect to visual coordinates.

Given the goal of recovering the view transform, how can the general ideas about context help? One of the simplest constraints that can be supplied by context is the knowledge of a supporting surface, the simplest of which is a plane. The viewing transformation has six parameters in general, but for most objects, the constraint of planar support reduces the degrees of freedom to three [71]. This is because most objects have very limited ways in which they can be supported by a plane. A normal kind of coffee mug (with a handle) will have four: right side up, upside down, and two ways of lying on its side. If we look at "mug ethology", the mug spends almost all of its time in the first position. This means that to find a coffee mug on a table, an overwhelmingly good bet is that it will be in one support relation with three degrees of freedom: two translation and one rotation. Since the degrees of freedom are the same for those of a two-dimensional planar problem, one might suspect that a pose computation is possible using only the two-dimensional image as advocated by Lowe [30, 31]. In fact this is possible and the mathematical form of these constraints is developed in [66]. This use of spatial information has emphasized the WHERE task of locating known objects. Just as important, but given short emphasis here, is the use of geometric cues in the WHAT task of object identification. Interestingly enough, much recent work in identification finds ways around computing pose directly, e.g. [26, 30, 31], by using features which are relatively view invariant.

5. Coordinated behaviors

The fixation frame with its small fovea allows the animate system to simplify its access to the environment. The idea is that, at any given instant only a relatively small number of features of the external world are registered but through perceptual actions the system can actively control the features that are

registered. A consequence of this ability is that with animate vision or more generally animate perception, systems can be built which learn to operate in a complex task domain without the associated explosion in the input feature vector required to represent all the elements of the domain.

Steven Whitehead has applied reinforcement learning ideas to the study of animate vision [63, 64]. Whitehead has been studying block stacking tasks. On each trial, the system is presented with a pile of colored blocks. A pile can consist of any number of blocks and they can be arranged in any configuration. Each block is uniformly colored and can be either red, green, or blue. The system can manipulate the pile by picking and placing objects. An object can be picked up only if its top is clear, and an object can be placed on another object only if the target object's top is clear. When the system arranges the blocks into a *successful configuration*, it receives a positive reward and the trial ends. A successful configuration is some predefined set of states which represents a desired outcome. For example, one simple block stacking task is for the system to learn to pick up a green block. In this case, the successful configurations consist just of those states where the system is holding a green object. The objective of the system is to learn algorithms for arranging arbitrary configurations of blocks into successful configurations.

Most reinforcement learning systems have static sensory systems. That is, the semantics of the feature vectors that describe the external state are defined a priori. Further, the input vector is defined so that each state is "sufficiently discriminable." Unfortunately as the complexity of the task domain increases, in particular as the number of "possibly relevant" objects in the task grows, the size of the static input vector (state representation) grows very quickly even though the number of relevant objects remains small. The problem is that with a static input vector if an object may be relevant to the task then it *must* be represented internally.

In contrast to static systems, systems using animate vision can avoid the combinatorial explosion of absolute representations by using "indexical representations". The basic idea behind an indexical representation is that the system shouldn't attempt to maintain an accurate representation of every item in the universe, but instead should only register objects and aspects (features) that are relevant to the task at hand [1]. For the block stacking problem, instead of assigning an absolute symbolic name to each item in the universe, such as "BLOCK-44", the system only registers objects (and their features) according to the functional roles they play in solving the task, such as "THE-BLOCK-I-AM-FIXATING". Whitehead's system uses both a fixation frame and an *attention frame* as shown in Fig. 11. Details may be found in [64].

Over a number of trials the system can learn to solve particular tasks. Figure 12 shows the number of steps used to solve the problem of picking up a green block. The disadvantage of this approach is that, so far, there is no good way to generalize it. However, that should not obscure the many important

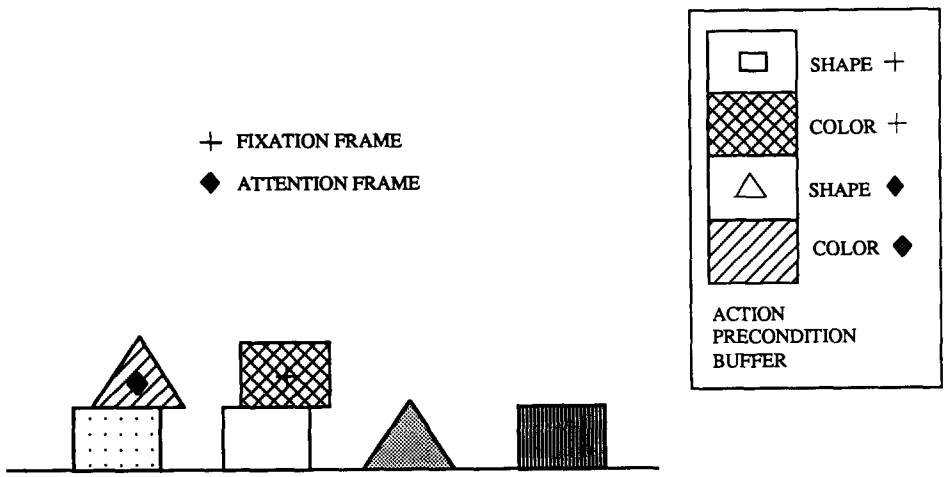

Fig. 11. The active perceptual system is divided into two parts: the fixation frame and an attention frame. The information registered in the fovea (or fixation point) can be actively controlled by executing perceptual and gaze control "acts". For example, fixating the block as shown causes its features to be registered in the state vector; attending to the triangle as shown causes its features to appear in the state vector. The two degrees of freedom in the state vector that can be independently controlled correspond to "markers". One can think of placing a special marker on an object causing its properties to appear in the appropriate place in the state vector. The system used by Whitehead is slightly more complex but still only uses twenty bits total to represent the state of the world.

conceptual points. We contend that searching huge state spaces such as those in blocks world domains may be impossible without the incorporation of these kinds of ideas in animate vision systems. First, learning by trial and error allows the agent to amortize building a policy function over its history. Once a good policy function is learned, applying it is cheap. Second, the reinforcement learning algorithm we use has a limited attention span, so that it gives up after expending a predetermined amount of resources. This is important because (a) a real-time system *has* to respond in a timely manner and (b) this strategy, in the context of repeated applications, causes the agent to gradually improve its competence [11]. The third advantage of this kind of learning derives from the use of indexical representation. This allows (a) the access of items by property instead of by category, and (b) run-time indexing. *Access by property* is efficient in the following way. Consider the problem of hanging a picture where a nail has to be driven into a wall. We do not really need a hammer, but something that could serve as a hammer. *Plan access by category* forces the identification of image items, followed by a check to determine the appropriate properties. *Access by properties* short circuits this process. Also, the fact that these properties are determined by what is in the environment at the moment filters out the consideration of strategies that would require unavailable items.

One problem such systems will have is the well-known credit assignment problem. If the reinforcements change the problem of how to change the

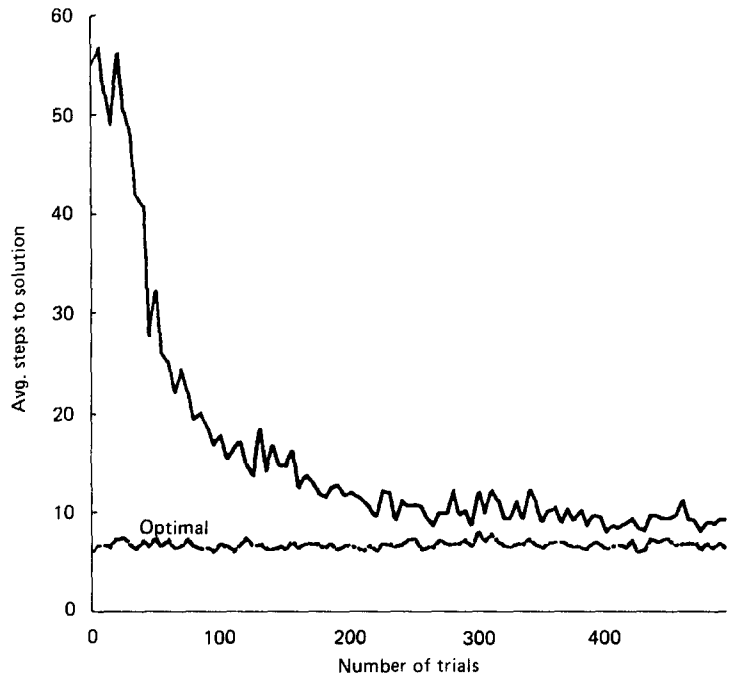

Fig. 12. The results of applying reinforcement learning to a simple block stacking task: Pick up the green block. The system uses its fixation and attention frames to register features dynamically and thus avoid the combinatorial cost of representing large state spaces. The lower trace shows the smallest number of steps needed to solve the problem computed as the running average of the last three presentations. The upper trace shows the time taken to solve the problem, also averaged over three presentations (from Whitehead and Ballard [64]).

reinforcement schedule is completely open. However, such a system can search local to the policy that it has fairly cheaply and this may work for an interesting set of behaviors.

6. Conclusions

An animate vision system with the ability to control its gaze can make the execution of behaviors involving vision much simpler. Gaze control confers several advantages in the use of vision in behavioral tasks, and these have been summarized in Section 2.

As humans we have the compelling experience of living in a three-dimensional visual movie. The world appears vividly colorful and stable. One temptation is to propose models of perception that capture this phenomenon in very explicit ways, say as a pictorial memory buffer. If the explicit buffer seems too crude, one can posit elaborate data structures that are equivalent in the sense that they contain the information necessary to construct such a picture.

However, when one examines the *mechanisms* of human and animal visual perception in detail, or tries to build anthropomorphic robots, it quickly becomes apparent that the way the apparatus works at this level of abstraction, e.g., the fast sequential saccadic searches, is incompatible with phenomenological notions of invariance and stability. Models of the visual system that work are compartmentalized with inconsistent representations and specialized behaviors that compete for the resources of the system. In this milieux, animate vision has a huge run-time component. Vision depends on the world being sufficiently stable so that behaviors can be executed on demand. Perhaps it is this ability to conduct behaviors that make assumptions about the world that provides the illusion of stable perception. Another way to say this is that: *Animate systems that rapidly change their coupling with the real world place a premium on maintaining elaborate representations of the world. However, it may be the case that memorizing such representations is unnecessary, since they can be rapidly and incrementally computed on demand.*

The ability to have behaviors that learn to adapt to the local environment will have a profound effect on the design of animate vision algorithms. The discussion on color introduced the notion of a personalized representation: that is, associating features with an object that makes the behaviors concerning it especially easy to execute. One can think of many other cases that challenge traditional notions of invariance. For example, we do not think of our coats as being rigid objects, yet they appear to be to our visual systems while they are hanging on coat racks, and this limited invariance can be exploited. The hope is that such algorithms may be able to discover which combinations of such features work in each problem instance. It could be the case that the general assumptions that define categories are almost never as useful as the special assumptions found by adaptive algorithms.

The study of animate vision is in its infancy, but we can already project that this paradigm will extend the capabilities of all kinds of computer vision systems, but particularly those of mobile vision platforms.

Acknowledgement

I would like to thank my colleagues Christopher Brown and Randal Nelson for critiquing these ideas. I am also especially grateful to the University of Rochester team of researchers, Tim Becker, David Coombs, Nat Martin, Tom Olson, Robert Potter, Ray Rimey, Michael Swain, Dave Tilley, Steve Whitehead, Lambert Wixson, and Brian Yamauchi, all of whom have greatly helped refine the ideas herein. The color examples are derived from research projects headed by Michael Swain and Lambert Wixson. The learning example comes from a research project headed by Steve Whitehead. Peggy Meeker is responsible for the pleasing format of the manuscript.

References

[1] P.E. Agre and D. Chapman, Pengi: an implementation of a theory of activity, in: *Proceedings AAAI-87*, Seattle, WA (1987) 268–272.

[2] J. Aloimonos, A. Bandopadhay and I. Weiss, Active vision, in: *Proceedings First International Conference on Computer Vision*, London (1987) 35–54; also *Int. J. Comput. Vision* **1** (4) (1988) 333–356.

[3] M.A. Arbib, Perceptual structures and distributed motor control, in: V.B. Brooks, ed., *Handbook of Physiology: The Nervous System II. Motor Control* (American Physiological Society, Bethesda, MD, 1981) 1449–1480.

[4] R. Bajcsy and P. Allen, Sensing strategies, in: *Proceedings U.S.-France Robotics Workshop*, Philadelphia, PA (1984).

[5] D.H. Ballard, Generalizing the Hough transform to arbitrary shapes, in: *Proceedings International Conference on Computer Vision and Pattern Recognition* (1981).

[6] D.H. Ballard, Behavioral constraints on computer vision, *Image Vision Comput.* **7** (1) (1989).

[7] D.H. Ballard and A. Ozcandarli, Eye fixation and early vision: kinetic depth, in: *Proceedings 2nd IEEE International Conference on Computer Vision*, Tampa, FL (1988).

[8] A. Bandopadhay, A computational study of rigid motion, Ph.D. Thesis, Computer Science Department, University of Rochester, Rochester, NY (1987).

[9] J.R. Beveridge, J. Griffith, R.R. Kohler, A.R. Hanson and E.M. Riseman, Segmenting images using localized histograms and region merging, *Int. J. Comput. Vision* **2** (3) (1989) 311–347.

[10] I. Biederman, Human image understanding: Recent research and a theory, *Comput. Vision Graph. Image Process.* **32** (1) (1985).

[11] L. Blum and M. Blum, Toward a mathematical theory of inductive inference, *Inf. Control* **22** (1975) 125–155.

[12] R.A. Brooks, Achieving artificial intelligence through building robots, TR 899, MIT, Cambridge, MA (1986).

[13] R.A. Brooks, A robust layered control system for a mobile robot, *IEEE J. Rob. Autom.* **2** (1986) 14–23.

[14] C.M. Brown, Gaze controls with interactions and delays, *IEEE Trans. Syst. Man Cybern.* **20** (2) (1990).

[15] C.M. Brown, Prediction and cooperation in gaze control, *Biol. Cybern.* (1990).

[16] C.M. Brown, ed., with D.H. Ballard, T.G. Becker, R.F. Gans, N.G. Martin, T.J. Olson, R.D. Potter, R.D. Rimey, D.G. Tilley and S.D. Whitehead, The Rochester robot, TR 257, Computer Science Department, University of Rochester, Rochester, NY (1988).

[17] C.H. Chen and A.C. Kak, A robot vision system for recognizing 3-d objects in low-order polynomial time, *IEEE Trans. Syst. Man Cybern.* (1989) Special Issue on Computer Vision.

[18] J.J. Clark and N.J. Ferrier, Modal control of an attentive vision system, in: *Proceedings 2nd International Conference on Computer Vision*, Tampa, FL (1988) 514–523.

[19] J.E. Cutting, Motion parallax and visual flow: How to determine direction of locomotion, in: *Proceedings Meeting of the International Society for Ecological Psychology*, Hartford, CT (1982).

[20] E.D. Dickmanns, Real-time machine vision exploiting integral spatio-temporal world models, invited presentation, *IJCAI-89*, Detroit, MI (1989).

[21] C.J. Erkelens and H. Collewijn, Eye movements and stereopsis during dichoptic viewing of moving random-dot stereograms, *Vision Res.* **25** (1985) 1689–1700.

[22] J.A. Feldman and Y. Yakimovsky, Decision theory and artificial intelligence I: A semantics-based region analyzer, *Artif. Intell.* **5** (1974) 349–371.

[23] T.D. Garvey, Perceptual strategies for purposive vision, Tech. Note 117, SRI International, Menlo Park, CA (1976).

[24] J.J. Gibson, *The Ecological Approach to Visual Perception* (Houghton Mifflin, Boston, MA, 1979).

[25] G.E. Hinton, Shape representation in parallel systems, in: *Proceedings IJCAI-81*, Vancouver, BC (1981).

[26] D.P. Huttenlocher and S. Ullman, Recognizing solid objects by alignment, in: *Proceedings DARPA Image Understanding Workshop* (1988).
[27] G.J. Klinker, S.A. Shafer and T. Kanade, The measurement of highlights in color images, *Int. J. Comput. Vision* **2** (1988) 7–32.
[28] H. Komatsu and R.H. Wurtz, Relation of cortical areas MT and MST to pursuit eye movements III: interaction with full-field visual stimulation, *J. Neurophysiol.* **60** (2) (1988) 621–644.
[29] E. Krotkov, Focusing, *Int. J. Comput. Vision* **1** (3) (1988) 223–238.
[30] D. Lowe, *Perceptual Organization and Visual Recognition* (Kluwer Academic Publishers, Boston, MA, 1985).
[31] D. Lowe, Fitting parameterized 3-d models to images, Tech. Rept. 89-26, Computer Science Department, University of British Columbia, Vancouver, BC (1989).
[32] L.T. Maloney and B.A. Wandell, Color constancy: A method for recovering surface spectral reflectance, *J. Optical Soc. Am. A* **3** (1) (1986) 29–33.
[33] D.C. Marr, *Vision* (Freeman, San Francisco, CA, 1982).
[34] J.H.R. Maunsell and W.T. Newsome, Visual processing in monkey extrastriate cortex, *Annu. Rev. Neurosci.* **10** (1987) 363–401.
[35] J.H.R. Maunsell and D. Van Essen, The topographic organization of the middle temporal visual area in the macaque monkey: representational biases and the relationship to callosal connections and mylo-architectonic boundaries, *J. Comparative Neurol.* **266** (1986) 535–555.
[36] M. Mishkin, A memory system in the monkey, *Philos. Trans. Royal Soc. London B* **298** (1982) 85–95.
[37] M. Mishkin, L.G. Ungerleider and K.A. Macko, Object vision and spatial vision: two cortical pathways, *Trends Neurosci.* **6** (1983) 414–417.
[38] H.P. Moravec, Towards automatic visual obstacle avoidance, in: *Proceedings IJCAI-77*, Cambridge, MA (1977) 584.
[39] K. Nakayama, Presentation, *Workshop on Computational Neuroscience*, Woods Hole, MA (1988).
[40] R.C. Nelson, Detection of motion by a moving observer: two qualitative approaches (submitted 1990).
[41] R.C. Nelson and J. Aloimonos, Obstacle avoidance using flow field divergence, *IEEE Trans. Pattern Anal. Mach. Intell.* **11** (1989) 1102–1106.
[42] F.W. Newell, *Ophthalmology: Principles and Concepts* (C.V. Mosby, St. Louis, MO, 1982).
[43] D. Noton, A theory of visual pattern perception, *IEEE Trans. Syst. Sci. Cybern.* **6** (1970) 349–357.
[44] D. Noton and L. Stark, Eye movements and visual perception, *Sci. Am.* **224** (6) (1971) 34–43.
[45] D. Noton and L. Stark, Scanpaths in saccadic eye movements while viewing and recognizing patterns, *Vision Res.* **11** (1971) 929.
[46] T.J. Olson and D.J. Coombs, Real-time vergence control for binocular robots, TR 348, Computer Science Department, University of Rochester, Rochester, NY (1990).
[47] T.J. Olson and R.D. Potter, Real-time vergence control, TR 264, Computer Science Department, University of Rochester, Rochester, NY (1988).
[48] J.K. O'Regan and A. Lévy-Schoen, Integrating visual information from successive fixations: does trans-saccadic fusion exist? *Vision Res.* **23** (8) (1983) 765–768.
[49] A. Pentland, A new sense of depth of field, in: *Proceedings IJCAI-85*, Los Angeles, CA (1985) 988–994.
[50] A. Pentland, Shape from shading: a theory of human perception, in: *Proceedings 2nd International Conference on Computer Vision*, Tampa, FL (1988).
[51] V.S. Ramachandran, Interactions between motion, depth, color and form: the utilitarian theory of perception, in: *Proceedings Conference on Visual Coding and Efficiency* (1987).
[52] D. Raviv and M. Herman, Towards an understanding of camera fixation, Tech. Rept., Robot Systems Division, National Institutes of Standards and Technology (1989).
[53] R.D. Rimey and C.M. Brown, Selective attention as sequential behavior: modeling eye movements with an augmented hidden Markov model, TR 327, Computer Science Department, University of Rochester, Rochester, NY (1990).

[54] S.J. Rosenschein, Formal theories of knowledge in AI and robotics, Tech. Note 362, AI Center, SRI International, Menlo Park, CA (1985).
[55] S.J. Rosenschein and L. Kaelbling, The synthesis of digital machines with provable epistemic properties, in: *Proceedings Conference on Theoretical Aspects of Reasoning about Knowledge*, Monterey, CA (1986).
[56] J. Rubner and K. Schulten, A regularized approach to color constancy, *Biol. Cybern.* **61** (1) (1989) 29–36.
[57] G. Sandini and V. Tagliasco, An anthropomorphic retina-like structure for scene analysis, *Comput. Vision Graph. Image Process.* **14** (4) (1980) 365–372.
[58] S. Shapiro, ed., *Encyclopedia of Artificial Intelligence* (Wiley, New York, 1987).
[59] M.J. Swain, Color indexing, Ph.D. Thesis, Computer Science Department, University of Rochester, NY (1990).
[60] M.J. Swain and D.H. Ballard, Object identification using color cues, Tech. Rept., Computer Science Department, University of Rochester, Rochester, NY (1990).
[61] J. Tsotsos, A complexity level analysis of vision, in: *Proceedings First International Conference on Computer Vision*, London (1987).
[62] S. Ullman, Visual routines, *Cognition* **18** (1984) 97–157; also in: S. Pinker, ed., *Visual Cognition* (MIT Press/Bradford Books, Cambridge, MA, 1984) 97–160.
[63] S.D. Whitehead and D.H. Ballard, A role for anticipation in reactive systems that learn, in: *Proceedings Sixth International Workshop on Machine Learning*, Ithaca, NY (1989).
[64] S.D. Whitehead and D.H. Ballard, Active perception and reinforcement learning (submitted 1990).
[65] L.E. Wixson and D.H. Ballard, Color histograms for real-time object search, in: *Proceedings SPIE Sensor Fusion II: Human and Machine Strategies Workshop*, Philadelphia, PA (1989).
[66] L.E. Wixson and D.H. Ballard, Detecting object pose using context, Tech. Rept., Computer Science Department, University of Rochester, Rochester, NY (1990).
[67] B. Yamauchi, JUGGLER: Real-time sensorimotor control using independent agents, in: *Proceedings Optical Society of America Image Understanding and Machine Vision Conference*, N. Falmouth, MA (1989) 6–9.
[68] A.L. Yarbus, *Eye Movements and Vision* (Plenum, New York, 1967).
[69] Y. Yeshurun and E.L. Schwartz, Cepstral filtering on a columnar image architecture: A fast algorithm for binocular stereo segmentation, Robotics Res. Tech. Rept. 286, Courant Institute, New York (1987).
[70] J.R. Lishman, Vision and the optic from field, *Nature* **293** (1981) 263–264.
[71] T.M. Siberberg, D.A. Harwood and L.S. Davis, Object recognition using oriented model points, *Comput. Vision Graph. Image Process.* **35** (1986) 47–71.